

MAIN

October 22, 2020

```
[19]: import seaborn as sns
import os
import pandas as pd
import numpy as np
import json
import plotly.graph_objects as go
import copy

from plotly.subplots import make_subplots
from scipy import stats

from misc import create_all_paths, get_class_dist_timeseries
from cleaning import load_arff_files
from feature_selecting import GreedyForwardSelector, CorrelationSelector
from modelling import CrossValidation, XGBoost, LogisticRegression
from preprocessing import Standardizer, MeanReplacement, clip_outliers, \
    ↳ ReSampler, CorrelationRemover, PCA, remove_outliers
from evaluation import get_f1_score, get_accuracy, get_auc, get_precision, \
    ↳ get_recall, get_all_measures, \
        get_threshold_for_optim_cost, get_weighted_accuracy, \
    ↳ get_params_for_best_measure_overall, get_auc, \
        get_params_for_best_model
from visualization import get_roc_curve
from __main_fs__ import main_corr, main_greedy, pred_function

# I am aware of this but get warnings nevertheless since i copy the dataframe in, \
↳ some steps which sets the attribute is_copy
pd.set_option('chained_assignment', None)
create_all_paths()

[3]: ##### Exploration #####
dfs, file_names = load_arff_files("../data/raw_data/")
assert len(dfs) > 0, "You forgot to copy the arff files inside the data/
↳ raw_data directory"
year_one = dfs[0]
year_one.head()
```

```

[3]: net profit / total assets    total liabilities / total assets    \
0          0.089414                0.656650
1          0.092379                0.643260
2         -1.372700                1.990500
3          0.056093                0.524820
4          0.084977                0.096562

    working capital / total assets    current assets / short-term liabilities    \
0          0.18322                1.33280
1          0.12866                1.26180
2         -1.06240                0.46628
3          0.30847                1.60930
4          0.32441                4.81130

    [(cash + short-term securities + receivables - short-term liabilities) /
    (operating expenses - depreciation)] * 365    \
0          -47.9670
1          22.0960
2          -7.8927
3         -26.7190
4          63.1220

    retained earnings / total assets    EBIT / total assets    \
0          0.00000                0.112660
1          0.00000                0.110500
2          0.00000                -1.372700
3          0.35300                0.071484
4          0.26805                0.085542

    book value of equity / total liabilities    sales / total assets    \
0          0.52288                2.2558
1          0.55457                2.2507
2         -0.49762                48.0050
3          0.88152                1.0442
4          9.29990                1.0984

    equity / total assets    ...    (sales - cost of products sold) / sales    \
0          0.34335    ...                0.280410
1          0.35674    ...                0.037282
2         -0.99052    ...               -0.009713
3          0.46264    ...                0.042332
4          0.89802    ...                0.089598

    (current assets - inventory - short-term liabilities) / (sales - gross profit
    - depreciation)    \
0          0.260420
1          0.258950

```

2	1.385900
3	0.121250
4	0.094627

	total costs /total sales	long-term liabilities / equity \
0	0.74418	0.225860
1	0.95226	0.221070
2	1.00730	0.000000
3	0.95767	0.040020
4	0.91040	0.012744

	sales / inventory	sales / receivables \
0	6.7518	6.7353
1	NaN	4.6606
2	NaN	63.7950
3	4.3271	6.4310
4	5.3233	9.4997

	(short-term liabilities *365) / sales	sales / short-term liabilities \
0	89.083	4.0973
1	79.689	4.5803
2	15.135	24.1170
3	96.845	3.7689
4	32.600	11.1960

	sales / fixed assets	class
0	8.4735	0
1	5.9233	0
2	668.0900	1
3	10.3020	0
4	1.6140	0

[5 rows x 65 columns]

```
[4]: year_one.describe()
# We have all numerical columns. No categorical/character columns.
# There seems to be some Missing values present. Lets observe them
```

```
[4]:
```

	net profit / total assets	total liabilities / total assets \
count	7024.000000	7024.000000
mean	0.034660	0.560215
std	4.565504	5.350084
min	-256.890000	-72.162000
25%	0.021182	0.296678
50%	0.075802	0.482960
75%	0.160268	0.680233
max	94.280000	441.500000

	working capital / total assets \
count	7024.000000
mean	0.119969
std	5.275459
min	-440.500000
25%	0.026968
50%	0.181275
75%	0.362548
max	1.000000

	current assets / short-term liabilities \
count	6997.000000
mean	2.629143
std	13.257356
min	0.000000
25%	1.063100
50%	1.502000
75%	2.460700
max	1017.800000

	[(cash + short-term securities + receivables - short-term liabilities) / (operating expenses - depreciation)] * 365 \
count	7.019000e+03
mean	-2.631672e+02
std	3.707460e+04
min	-2.722100e+06
25%	-4.449800e+01
50%	-5.373900e+00
75%	3.777050e+01
max	9.909000e+05

	retained earnings / total assets	EBIT / total assets \
count	7024.000000	7024.000000
mean	0.059712	0.313876
std	6.051113	8.353274
min	-397.890000	-189.560000
25%	0.000000	0.028023
50%	0.000000	0.090109
75%	0.146660	0.188667
max	303.670000	453.770000

	book value of equity / total liabilities	sales / total assets \
count	7002.000000	7026.000000
mean	2.623996	5.552855
std	18.708327	101.995448
min	-141.410000	0.000005

25%	0.445710	1.037225
50%	1.015100	1.205750
75%	2.267675	2.132975
max	1452.200000	3876.100000

	equity / total assets ...	(sales - cost of products sold) / sales \
count	7024.000000 ...	7.027000e+03
mean	1.825832 ...	-1.577367e+02
std	33.836452 ...	1.322125e+04
min	-440.550000 ...	-1.108300e+06
25%	0.300785 ...	2.031450e-02
50%	0.492235 ...	6.338200e-02
75%	0.675677 ...	1.376950e-01
max	1099.500000 ...	1.000000e+00

	(current assets - inventory - short-term liabilities) / (sales - gross profit - depreciation) \
count	7026.000000
mean	0.193243
std	4.344046
min	-315.370000
25%	0.056772
50%	0.175745
75%	0.351922
max	126.670000

	total costs /total sales	long-term liabilities / equity \
count	7.027000e+03	7026.000000
mean	1.587409e+02	0.277829
std	1.322124e+04	6.339149
min	-4.194000e-03	-327.970000
25%	8.647650e-01	0.000000
50%	9.388100e-01	0.028438
75%	9.820150e-01	0.273867
max	1.108300e+06	119.580000

	sales / inventory	sales / receivables \
count	6.892000e+03	7005.000000
mean	4.328830e+02	15.642228
std	2.612802e+04	261.554534
min	4.700000e-05	0.000016
25%	5.923950e+00	4.829000
50%	1.004050e+01	7.033700
75%	2.013900e+01	10.703000
max	2.137800e+06	21110.000000

(short-term liabilities *365) / sales \

count	7.027000e+03
mean	4.763202e+03
std	3.107835e+05
min	0.000000e+00
25%	4.322250e+01
50%	6.850900e+01
75%	1.063350e+02
max	2.501600e+07

	sales / short-term liabilities	sales / fixed assets	class
count	6997.000000	6993.000000	7027.000000
mean	8.126852	208.731950	0.038566
std	19.996419	5140.708804	0.192571
min	0.000015	0.000010	0.000000
25%	3.425400	2.538600	0.000000
50%	5.303200	4.637700	0.000000
75%	8.357900	9.782200	0.000000
max	1042.200000	294770.000000	1.000000

[8 rows x 65 columns]

```
[5]: number_nas = (~np.isfinite(year_one)).sum(axis=0)
print("NAs in Total:{} %".format(round(sum(number_nas)/(year_one.
    ↳shape[1]*year_one.shape[0])*100, 2)))
for na, col in zip(number_nas, year_one.columns):
    print("{} / {} % missing values in column {}".format(na, round(na/year_one.
    ↳shape[0]*100, 2), col[:60]))
```

```
NAs in Total:1.28 %
3 / 0.04 % missing values in column net profit / total assets
3 / 0.04 % missing values in column total liabilities / total assets
3 / 0.04 % missing values in column working capital / total assets
30 / 0.43 % missing values in column current assets / short-term liabilities
8 / 0.11 % missing values in column [(cash + short-term securities + receivables
- short-term li
3 / 0.04 % missing values in column retained earnings / total assets
3 / 0.04 % missing values in column EBIT / total assets
25 / 0.36 % missing values in column book value of equity / total liabilities
1 / 0.01 % missing values in column sales / total assets
3 / 0.04 % missing values in column equity / total assets
39 / 0.56 % missing values in column (gross profit + extraordinary items +
financial expenses) /
30 / 0.43 % missing values in column gross profit / short-term liabilities
0 / 0.0 % missing values in column (gross profit + depreciation) / sales
3 / 0.04 % missing values in column (gross profit + interest) / total assets
2 / 0.03 % missing values in column (total liabilities * 365) / (gross profit +
depreciation)
25 / 0.36 % missing values in column (gross profit + depreciation) / total
```

liabilities

25 / 0.36 % missing values in column total assets / total liabilities

3 / 0.04 % missing values in column gross profit / total assets

0 / 0.0 % missing values in column gross profit / sales

0 / 0.0 % missing values in column (inventory * 365) / sales

1622 / 23.08 % missing values in column sales (n) / sales (n-1)

3 / 0.04 % missing values in column profit on operating activities / total assets

0 / 0.0 % missing values in column net profit / sales

124 / 1.76 % missing values in column gross profit (in 3 years) / total assets

3 / 0.04 % missing values in column (equity - share capital) / total assets

25 / 0.36 % missing values in column (net profit + depreciation) / total liabilities

311 / 4.43 % missing values in column profit on operating activities / financial expenses

34 / 0.48 % missing values in column working capital / fixed assets

3 / 0.04 % missing values in column logarithm of total assets

0 / 0.0 % missing values in column (total liabilities - cash) / sales

0 / 0.0 % missing values in column (gross profit + interest) / sales

38 / 0.54 % missing values in column (current liabilities * 365) / cost of products sold

30 / 0.43 % missing values in column operating expenses / short-term liabilities

25 / 0.36 % missing values in column operating expenses / total liabilities

3 / 0.04 % missing values in column profit on sales / total assets

3 / 0.04 % missing values in column total sales / total assets

2740 / 38.99 % missing values in column (current assets - inventories) / long-term liabilities

3 / 0.04 % missing values in column constant capital / total assets

0 / 0.0 % missing values in column profit on sales / sales

30 / 0.43 % missing values in column (current assets - inventory - receivables) / short-term liab

84 / 1.2 % missing values in column total liabilities / ((profit on operating activities + depre

0 / 0.0 % missing values in column profit on operating activities / sales

0 / 0.0 % missing values in column rotation receivables + inventory turnover in days

0 / 0.0 % missing values in column (receivables * 365) / sales

134 / 1.91 % missing values in column net profit / inventory

31 / 0.44 % missing values in column (current assets - inventory) / short-term liabilities

29 / 0.41 % missing values in column (inventory * 365) / cost of products sold

3 / 0.04 % missing values in column EBITDA (profit on operating activities - depreciation) / tot

0 / 0.0 % missing values in column EBITDA (profit on operating activities - depreciation) / sal

25 / 0.36 % missing values in column current assets / total liabilities

3 / 0.04 % missing values in column short-term liabilities / total assets

29 / 0.41 % missing values in column (short-term liabilities * 365) / cost of

```

products sold)
34 / 0.48 % missing values in column equity / fixed assets
34 / 0.48 % missing values in column constant capital / fixed assets
0 / 0.0 % missing values in column working capital
0 / 0.0 % missing values in column (sales - cost of products sold) / sales
1 / 0.01 % missing values in column (current assets - inventory - short-term
liabilities) / (sal
0 / 0.0 % missing values in column total costs /total sales
1 / 0.01 % missing values in column long-term liabilities / equity
135 / 1.92 % missing values in column sales / inventory
22 / 0.31 % missing values in column sales / receivables
0 / 0.0 % missing values in column (short-term liabilities *365) / sales
30 / 0.43 % missing values in column sales / short-term liabilities
34 / 0.48 % missing values in column sales / fixed assets
0 / 0.0 % missing values in column class

```

```

[6]: correlations = year_one.corr()
fig = go.Figure(data=go.Heatmap(
    z=correlations,
    y=correlations.index,
    x=correlations.index))
fig.update_layout(xaxis={'showgrid': False, 'visible': False},
    yaxis={'showgrid': False, 'visible': False})
fig.show()

# For a quick solution to reduce the multicollinearity we should drop the
→following columns since they have the
# strongest and most correlations to other columns:
# EBITDA / sales, (sales - cost of products sold) / sales, total costs/total
→sales, receivables*365 / sales, rotation receivables + inventory turnover in
→days,
# profit on operating activities/sales

```

```

[7]: # As expected, we have an unbalanced dataset. The positive class has a small
→share with less than 7% of all data
# Therefore, I need to rebalance the dataset while training to prevent the
→algorithm to focus the negative class since it
# leads to a larger decrease of costs
for df, file_name in zip(dfs, file_names):
    print("Class Distribution for '{}' :".format(file_name))
    print(df["class"].value_counts())

```

```

Class Distribution for '../data/raw_data/1year.arff' :
0    6756
1     271
Name: class, dtype: int64
Class Distribution for '../data/raw_data/2year.arff' :
0    9773

```



```

1      400
Name: class, dtype: int64
Class Distribution for '../data/raw_data/3year.arff' :
0      10008
1         495
Name: class, dtype: int64
Class Distribution for '../data/raw_data/4year.arff' :
0       9277
1        515
Name: class, dtype: int64
Class Distribution for '../data/raw_data/5year.arff' :
0       5500
1        410
Name: class, dtype: int64

```

```

[8]: bankruptcy_timeseries = get_class_dist_timeseries([df["class"] for df in dfs])
fig = make_subplots(specs=[[{"secondary_y": True}]]))
fig.add_trace(go.Scatter(x=bankruptcy_timeseries["Year"],
    ↪y=bankruptcy_timeseries["Share of bankrupt customers"],
                        mode='lines', name="Share of bankrupt customers"),
    ↪secondary_y=True)
fig.add_trace(go.Scatter(x=bankruptcy_timeseries["Year"],
    ↪y=bankruptcy_timeseries["Customers"],
                        mode='lines', name='Number of Customers'),
    ↪secondary_y=False)
fig.update_layout(paper_bgcolor='rgba(0,0,0,0)', plot_bgcolor='rgba(0,0,0,0)')

```

```

[9]: rows, cols = 18, 4
columns = year_one.columns
fig = make_subplots(rows=rows, cols=cols, subplot_titles=["{} - {}".format(i,
    ↪col[:15]) for i, col in enumerate(columns)])
for i in range(year_one.shape[1]):
    row_val = int(i/cols) + 1
    col_val = (i) % cols + 1
    fig.add_trace(go.Histogram(x=remove_outliers(year_one.iloc[:, i], 3.5),
    ↪name="{} - {}".format(i, columns[i][:15])),
                    row=row_val, col=col_val)
fig.update_layout(height=1600, width=1000)
fig.show()
# Logtransform the following columns:
# 3, 7, 16, 19, 31, 32, 33, 35, 36, 39, 42, 43, 45, 46, 49, 50, 51, 59, 60-63
# invert and root squar the following columns:
# 9, 24, 37

```

```

[10]: go.Figure(go.Histogram(x=remove_outliers((year_one.iloc[:, 3]), 5.5))).show()
go.Figure(go.Histogram(x=remove_outliers(np.log(year_one.iloc[:, 3]), 5.5))).
    ↪show()

```

c:\users\felix\appdata\local\programs\python\python37\lib\site-packages\pandas\core\series.py:726: RuntimeWarning:

divide by zero encountered in log

```
[11]: ##### PCA #####
tmp_pca = dfs[0].copy()
tmp_pca = tmp_pca.drop("class", axis=1)
pca = PCA(tmp_pca.shape[1])
tmp_pca = MeanReplacement().process(tmp_pca)[0]
tmp_pca = Standardizer().process(tmp_pca)[0]

_ = pca.process(tmp_pca)
variance_ratio = pca._pca_model.explained_variance_ratio_
cum_ratio = np.cumsum(variance_ratio)

fig = go.Figure()
fig.add_trace(go.Scatter(x=np.arange(len(variance_ratio)), y=cum_ratio,
                        mode='lines', name='Cumulated relative variance'))
fig.update_layout(paper_bgcolor='rgba(0,0,0,0)', plot_bgcolor='rgba(0,0,0,0)',
                  title="PCA - Cumulated relative Variance")
fig.update_layout(shapes=[
    dict(type= 'line', yref= 'paper', y0= 0, y1= 1, xref= 'x', x0= 10, x1= 10),
    dict(type= 'line', yref= 'paper', y0= 0, y1= 1, xref= 'x', x0= 16, x1= 16),
    dict(type= 'line', yref= 'paper', y0= 0, y1= 1, xref= 'x', x0= 20, x1= 20),
    dict(type= 'line', yref= 'paper', y0= 0, y1= 1, xref= 'x', x0= 25, x1= 25)
])
fig.show()
```

```
[12]: ##### MODELLING #####
cost_weight = 20
primary_measure = "Weighted Accuracy"

probs_lr = []
params = {"model": {"max_iter": 1000}, "preprocessors": [MeanReplacement(),
    ↳Standardizer(), PCA(0.999),
                                                ReSampler("down")]}

dfs, file_names = load_arff_files("../data/raw_data/")

cv = CrossValidation(folds=10)
all_measures = {}
if os.path.exists("../data/cleaned_data/recorded_measures_lr.json"):
    with open("../data/cleaned_data/recorded_measures_lr.json", "r") as f:
        recorded_measures = json.load(f)
```

```

else:
    recorded_measures = None

for i, df in enumerate(dfs):
    print("#### Year {}".format(i + 1))
    df = df.copy()
    label = df.pop("class")
    # this is to kick only to kick most extreme outlier
    df = clip_outliers(df.copy(), 5.5)
    model = LogisticRegression(max_iter=params["model"]["max_iter"])
    probs = cv.run(data=df, label=label, model=model,
                   preprocessors=params["preprocessors"])
    probs_lr.append(probs) #
    measures = get_all_measures(probs, label, 0.5)

    base_line_measures = get_all_measures(np.random.choice([0, 1],
    ↪size=len(label)), label, 0.5)
    print("F1, Baseline: {}, LogisticRegression: {}".
    ↪format(base_line_measures["f1"], measures["f1"]))
    print("Acc, Baseline: {}, LogisticRegression: {}".
    ↪format(base_line_measures["acc"], measures["acc"]))
    print("AUC, Baseline: {}, LogisticRegression: {}".
    ↪format(base_line_measures["auc"], measures["auc"]))
    print("Recall, Baseline: {}, LogisticRegression: {}".
    ↪format(base_line_measures["recall"], measures["recall"]))
    print("Precision, Baseline: {}, LogisticRegression: {}".
    ↪format(base_line_measures["precision"],
    ↪measures["precision"])))

    threshold, costs = get_threshold_for_optim_cost(probs, label,
    ↪weight=cost_weight)

    print("Threshold: {}".format(threshold))
    print("Costs: {}".format(costs))
    weighted_accuracy = get_weighted_accuracy([1 if p > threshold else 0 for p
    ↪in probs], label, cost_weight)
    measures["Weighted Accuracy"] = weighted_accuracy
    measures["params"] = copy.deepcopy(params)
    measures["params"]["preprocessors"] = [(p.__class__.__name__, p.__dict__)
    ↪for p in params["preprocessors"]]
    measures["probs"] = probs.tolist()
    measures["label"] = label.tolist()
    print("Weighted Accuracy, Baseline: {}, LogisticRegression: {}".format(
    ↪get_weighted_accuracy(np.random.choice([0, 1], len(label)), label,
    ↪cost_weight),

```

```

        weighted_accuracy
    ))
    print("All measures with cost optimal threshold: {}".
    →format(get_all_measures(probs, label, threshold)))
    all_measures["LR Year {}".format(i + 1)] = measures
    if recorded_measures is not None:
        if measures[primary_measure] > np.max([m["LR Year {}".format(i +
    →1)] [primary_measure]
                                                for m in recorded_measures]):
            print("Found new best measure {} with value {}".
    →format(primary_measure, measures[primary_measure]))

if recorded_measures is None:
    recorded_measures = [all_measures]
else:
    recorded_measures.append(all_measures)
with open("../data/cleaned_data/recorded_measures_lr.json", "w") as f:
    json.dump(recorded_measures, f)

get_roc_curve(probs_lr[i], label)

```

```

##### 1 Year #####
F1, Baseline: 0.06481965499215893, LogisticRegression: 0.1573926868044515
Acc, Baseline: 0.4908211185427636, LogisticRegression: 0.6983065319481998
AUC, Baseline: 0.47485984850967516, LogisticRegression: 0.7761312071380038
Recall, Baseline: 0.4575645756457565, LogisticRegression: 0.7306273062730627
Precision, Baseline: 0.03488045007032349, LogisticRegression:
0.08819599109131403
Threshold: 0.536
Costs: 3429
Weighted Accuracy, Baseline: 0.4963042049934297, LogisticRegression:
0.7183804204993429
All measures with cost optimal threshold: {'auc': 0.7761312071380038, 'f1':
0.16876122082585276, 'acc': 0.7364451401736161, 'recall': 0.6937269372693727,
'precision': 0.09606540623403168}
Found new best measure Weighted Accuracy with value 0.7183804204993429
##### 2 Year #####

```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-12-2babe9752394> in <module>
    25     model = LogisticRegression(max_iter=params["model"]["max_iter"])
    26     probs = cv.run(data=df, label=label, model=model,
---> 27         preprocessors=params["preprocessors"])

    28     probs_lr.append(probs) #
    29     measures = get_all_measures(probs, label, 0.5)

```

```

D:\CodingProjects\AccentureApplication\src\modelling.py in run(self, data,
↳ label, model, preprocessors)
    128         predictions = np.array([np.NaN for _ in range(len(indices))])
    129         for i in range(self.folds):
--> 130             train_data, test_data = data.iloc[[k for k in range(data.
↳ shape[0]) if k not in test_indices[i]], :], \
    131                                     data.iloc[test_indices[i], :])
    132             train_label, test_label = [v for j, v in enumerate(label) if
↳ j not in test_indices[i]], \

D:\CodingProjects\AccentureApplication\src\modelling.py in <listcomp>(.0)
    128         predictions = np.array([np.NaN for _ in range(len(indices))])
    129         for i in range(self.folds):
--> 130             train_data, test_data = data.iloc[[k for k in range(data.
↳ shape[0]) if k not in test_indices[i]], :], \
    131                                     data.iloc[test_indices[i], :])
    132             train_label, test_label = [v for j, v in enumerate(label) if
↳ j not in test_indices[i]], \

KeyboardInterrupt:

```

```

[13]: probs_xgb = []
dfs, file_names = load_arff_files("../data/raw_data/")
params = {"model": dict(val_share=0.2, n_rounds=8, lambda_=5,
↳ additional_booster_params={"params": {"max_depth": 4,
↳ "subsample": 0.5,
↳ "colsample_bytree": 0.5}}},
↳ verbose=True), "preprocessors": []}

cv = CrossValidation(folds=10)
all_measures = {}
if os.path.exists("../data/cleaned_data/recorded_measures_xgb.json"):
    with open("../data/cleaned_data/recorded_measures_xgb.json", "r") as f:
        recorded_measures = json.load(f)
else:
    recorded_measures = None

for i, df in enumerate(dfs):
    print("##### {} Year #####".format(i + 1))
    df = df.copy()
    label = df.pop("class")
    model = XGBoost(**params["model"]) #, "scale_pos_weight": (df.shape[0] - np.
↳ sum(np.sum(df["class"]))) / np.sum(df["class"]))
    probs = cv.run(data=df, label=label, model=model,
↳ preprocessors=params["preprocessors"])

```

```

probs_xgb.append(probs)

measures = get_all_measures(probs, label, 0.5)
base_line_measures = get_all_measures(np.random.choice([0, 1],
↪size=len(label)), label, 0.5)

print("F1, Baseline: {}, XGBoost: {}".format(base_line_measures["f1"],
↪measures["f1"]))
print("Acc, Baseline: {}, XGBoost: {}".format(get_accuracy([1 for _ in
↪range(len(label))], label),
measures["acc"]))
print("AUC, Baseline: {}, XGBoost: {}".format(base_line_measures["auc"],
↪measures["auc"]))
print("Recall, Baseline: {}, XGBoost: {}".
↪format(base_line_measures["recall"], measures["recall"]))
print("Precision, Baseline: {}, XGBoost: {}".
↪format(base_line_measures["precision"],
measures["precision"]))

threshold, costs = get_threshold_for_optim_cost(probs, label,
↪weight=cost_weight)
print("Threshold: {}".format(threshold))
print("Costs: {}".format(costs))
weighted_accuracy = get_weighted_accuracy([1 if p > threshold else 0 for p
↪in probs], label, cost_weight)
measures["Weighted Accuracy"] = weighted_accuracy
measures["params"] = copy.deepcopy(params)
measures["params"]["preprocessors"] = [(p.__class__.__name__, p.__dict__)
↪for p in params["preprocessors"]]
measures["probs"] = probs.tolist()
measures["label"] = label.tolist()
print("Weighted Accuracy, Baseline: {}, XGB: {}".format(
get_weighted_accuracy(np.random.choice([0, 1], len(label)), label,
↪cost_weight),
weighted_accuracy
))
print("All measures with cost optimal threshold: {}".
↪format(get_all_measures(probs, label, threshold)))
all_measures["XGB Year {}".format(i + 1)] = measures
if recorded_measures is not None:
    if measures[primary_measure] > np.max([m["XGB Year {}".format(i +
↪1)] [primary_measure]
for m in recorded_measures]):
        print("Found new best measure {} with value {}".
↪format(primary_measure, measures[primary_measure]))

if recorded_measures is None:

```

```

        recorded_measures = [all_measures]
    else:
        recorded_measures.append(all_measures)
    with open("../data/cleaned_data/recorded_measures_xgb.json", "w") as f:
        json.dump(recorded_measures, f)

get_roc_curve(probs_xgb[i], label)

```

```

##### 1 Year #####
[0]    train-logloss:0.47751    eval-logloss:0.47417
[1]    train-logloss:0.34625    eval-logloss:0.35450
[2]    train-logloss:0.26195    eval-logloss:0.28129
[3]    train-logloss:0.21269    eval-logloss:0.23499
[4]    train-logloss:0.18021    eval-logloss:0.20286
[5]    train-logloss:0.14959    eval-logloss:0.18266
[6]    train-logloss:0.13082    eval-logloss:0.17082
[7]    train-logloss:0.11753    eval-logloss:0.16327
[0]    train-logloss:0.47834    eval-logloss:0.47785
[1]    train-logloss:0.34669    eval-logloss:0.35716
[2]    train-logloss:0.26194    eval-logloss:0.28508
[3]    train-logloss:0.21393    eval-logloss:0.23770
[4]    train-logloss:0.18157    eval-logloss:0.20842
[5]    train-logloss:0.15140    eval-logloss:0.18881
[6]    train-logloss:0.13066    eval-logloss:0.17716
[7]    train-logloss:0.11994    eval-logloss:0.16892
[0]    train-logloss:0.47386    eval-logloss:0.47628
[1]    train-logloss:0.34424    eval-logloss:0.36247
[2]    train-logloss:0.25968    eval-logloss:0.29566
[3]    train-logloss:0.21156    eval-logloss:0.25096
[4]    train-logloss:0.17793    eval-logloss:0.22222
[5]    train-logloss:0.14806    eval-logloss:0.20543
[6]    train-logloss:0.12898    eval-logloss:0.19646
[7]    train-logloss:0.11791    eval-logloss:0.19000
[0]    train-logloss:0.47771    eval-logloss:0.47668
[1]    train-logloss:0.34867    eval-logloss:0.35782
[2]    train-logloss:0.26356    eval-logloss:0.28531
[3]    train-logloss:0.21505    eval-logloss:0.23684
[4]    train-logloss:0.18257    eval-logloss:0.20751
[5]    train-logloss:0.15341    eval-logloss:0.18885
[6]    train-logloss:0.13389    eval-logloss:0.17937
[7]    train-logloss:0.12308    eval-logloss:0.17136
[0]    train-logloss:0.47565    eval-logloss:0.47273
[1]    train-logloss:0.34539    eval-logloss:0.35226
[2]    train-logloss:0.25956    eval-logloss:0.27842
[3]    train-logloss:0.21102    eval-logloss:0.23120
[4]    train-logloss:0.17835    eval-logloss:0.19914
[5]    train-logloss:0.14882    eval-logloss:0.17924
[6]    train-logloss:0.12897    eval-logloss:0.16827

```

[7]	train-logloss:0.11692	eval-logloss:0.16062
[0]	train-logloss:0.47789	eval-logloss:0.47803
[1]	train-logloss:0.34791	eval-logloss:0.35996
[2]	train-logloss:0.26478	eval-logloss:0.28755
[3]	train-logloss:0.21592	eval-logloss:0.24113
[4]	train-logloss:0.18291	eval-logloss:0.20959
[5]	train-logloss:0.15437	eval-logloss:0.19060
[6]	train-logloss:0.13410	eval-logloss:0.18097
[7]	train-logloss:0.12193	eval-logloss:0.17303
[0]	train-logloss:0.47598	eval-logloss:0.47374
[1]	train-logloss:0.34517	eval-logloss:0.35296
[2]	train-logloss:0.26049	eval-logloss:0.27969
[3]	train-logloss:0.21310	eval-logloss:0.23075
[4]	train-logloss:0.17916	eval-logloss:0.19960
[5]	train-logloss:0.15044	eval-logloss:0.18010
[6]	train-logloss:0.13040	eval-logloss:0.16883
[7]	train-logloss:0.11856	eval-logloss:0.15813
[0]	train-logloss:0.47758	eval-logloss:0.47579
[1]	train-logloss:0.34852	eval-logloss:0.35725
[2]	train-logloss:0.26349	eval-logloss:0.28577
[3]	train-logloss:0.21570	eval-logloss:0.23822
[4]	train-logloss:0.18193	eval-logloss:0.20622
[5]	train-logloss:0.15227	eval-logloss:0.18926
[6]	train-logloss:0.13278	eval-logloss:0.17976
[7]	train-logloss:0.12127	eval-logloss:0.17119
[0]	train-logloss:0.47640	eval-logloss:0.47841
[1]	train-logloss:0.34579	eval-logloss:0.35899
[2]	train-logloss:0.25960	eval-logloss:0.28615
[3]	train-logloss:0.21131	eval-logloss:0.24085
[4]	train-logloss:0.17871	eval-logloss:0.20978
[5]	train-logloss:0.14947	eval-logloss:0.19138
[6]	train-logloss:0.12951	eval-logloss:0.18150
[7]	train-logloss:0.11825	eval-logloss:0.17442
[0]	train-logloss:0.47546	eval-logloss:0.47292
[1]	train-logloss:0.34423	eval-logloss:0.35461
[2]	train-logloss:0.25815	eval-logloss:0.28212
[3]	train-logloss:0.21002	eval-logloss:0.23508
[4]	train-logloss:0.17720	eval-logloss:0.20622
[5]	train-logloss:0.14703	eval-logloss:0.18803
[6]	train-logloss:0.12806	eval-logloss:0.17760
[7]	train-logloss:0.11684	eval-logloss:0.17027

F1, Baseline: 0.07391763463569166, XGBoost: 0.4431818181818182
 Acc, Baseline: 0.03856553294435748, XGBoost: 0.9721075850291732
 AUC, Baseline: 0.5083765913147587, XGBoost: 0.892825893179003
 Recall, Baseline: 0.5166051660516605, XGBoost: 0.2878228782287823
 Precision, Baseline: 0.0398066533977822, XGBoost: 0.9629629629629629
 Threshold: 0.102
 Costs: 1832

Weighted Accuracy, Baseline: 0.4929369250985545, XGB: 0.8495400788436268
All measures with cost optimal threshold: {'auc': 0.892825893179003, 'f1':
0.44469026548672563, 'acc': 0.9285612636971681, 'recall': 0.7416974169741697,
'precision': 0.3175355450236967}

2 Year

[0]	train-logloss:0.47752	eval-logloss:0.47723
[1]	train-logloss:0.35342	eval-logloss:0.36116
[2]	train-logloss:0.27544	eval-logloss:0.29049
[3]	train-logloss:0.22536	eval-logloss:0.24404
[4]	train-logloss:0.19322	eval-logloss:0.21257
[5]	train-logloss:0.16683	eval-logloss:0.19414
[6]	train-logloss:0.14558	eval-logloss:0.18422
[7]	train-logloss:0.13402	eval-logloss:0.17631
[0]	train-logloss:0.47853	eval-logloss:0.47505
[1]	train-logloss:0.35417	eval-logloss:0.35735
[2]	train-logloss:0.27598	eval-logloss:0.28435
[3]	train-logloss:0.22911	eval-logloss:0.23769
[4]	train-logloss:0.19789	eval-logloss:0.20747
[5]	train-logloss:0.17186	eval-logloss:0.18852
[6]	train-logloss:0.15119	eval-logloss:0.17713
[7]	train-logloss:0.14050	eval-logloss:0.16829
[0]	train-logloss:0.47572	eval-logloss:0.47592
[1]	train-logloss:0.35117	eval-logloss:0.36046
[2]	train-logloss:0.27388	eval-logloss:0.29187
[3]	train-logloss:0.22429	eval-logloss:0.24658
[4]	train-logloss:0.19327	eval-logloss:0.21833
[5]	train-logloss:0.16750	eval-logloss:0.19950
[6]	train-logloss:0.14741	eval-logloss:0.18965
[7]	train-logloss:0.13770	eval-logloss:0.18211
[0]	train-logloss:0.47747	eval-logloss:0.47845
[1]	train-logloss:0.35853	eval-logloss:0.36172
[2]	train-logloss:0.28527	eval-logloss:0.29020
[3]	train-logloss:0.23805	eval-logloss:0.24541
[4]	train-logloss:0.20499	eval-logloss:0.21637
[5]	train-logloss:0.18309	eval-logloss:0.19757
[6]	train-logloss:0.15640	eval-logloss:0.18687
[7]	train-logloss:0.14543	eval-logloss:0.18013
[0]	train-logloss:0.47951	eval-logloss:0.47906
[1]	train-logloss:0.36140	eval-logloss:0.36157
[2]	train-logloss:0.28947	eval-logloss:0.28989
[3]	train-logloss:0.24204	eval-logloss:0.24614
[4]	train-logloss:0.20871	eval-logloss:0.21655
[5]	train-logloss:0.18685	eval-logloss:0.19690
[6]	train-logloss:0.16512	eval-logloss:0.18562
[7]	train-logloss:0.15432	eval-logloss:0.17705
[0]	train-logloss:0.47568	eval-logloss:0.47692
[1]	train-logloss:0.35538	eval-logloss:0.35790
[2]	train-logloss:0.27866	eval-logloss:0.28493

```
[3]      train-logloss:0.23212      eval-logloss:0.24240
[4]      train-logloss:0.19805      eval-logloss:0.21472
[5]      train-logloss:0.17615      eval-logloss:0.19581
[6]      train-logloss:0.15520      eval-logloss:0.18283
[7]      train-logloss:0.14301      eval-logloss:0.17572
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-13-afe5c48075c3> in <module>
    19     label = df.pop("class")
    20     model = XGBoost(**params["model"]) #, "scale_pos_weight": (df.
↳shape[0] - np.sum(np.sum(df["class"]))) / np.sum(df["class"])
--> 21     probs = cv.run(data=df, label=label, model=model,
↳preprocessors=params["preprocessors"])
    22     probs_xgb.append(probs)
    23

D:\CodingProjects\AccentureApplication\src\modelling.py in run(self, data,
↳label, model, preprocessors)
    128         predictions = np.array([np.NaN for _ in range(len(indices))])
    129         for i in range(self.folds):
--> 130             train_data, test_data = data.iloc[[k for k in range(data.
↳shape[0]) if k not in test_indices[i]], :], \
    131                                     data.iloc[test_indices[i], :])
    132             train_label, test_label = [v for j, v in enumerate(label) i
↳j not in test_indices[i]], \

D:\CodingProjects\AccentureApplication\src\modelling.py in <listcomp>(.0)
    128         predictions = np.array([np.NaN for _ in range(len(indices))])
    129         for i in range(self.folds):
--> 130             train_data, test_data = data.iloc[[k for k in range(data.
↳shape[0]) if k not in test_indices[i]], :], \
    131                                     data.iloc[test_indices[i], :])
    132             train_label, test_label = [v for j, v in enumerate(label) i
↳j not in test_indices[i]], \

KeyboardInterrupt:
```

```
[15]: ##### EVALUATAION #####

params_xgb = get_params_for_best_measure_overall("auc", "../data/cleaned_data/
↳recorded_measures_xgb.json")
params_lr = get_params_for_best_measure_overall("auc", "../data/cleaned_data/
↳recorded_measures_lr.json")
```

```

model_first_year_xgb = get_params_for_best_model("auc", 1, "../data/
↳cleaned_data/recorded_measures_xgb.json")
model_first_year_lr = get_params_for_best_model("auc", 1, "../data/cleaned_data/
↳recorded_measures_lr.json")

measure = "auc"
measures = {"xgb": [params_xgb[key][measure] for key in params_xgb.keys()],
            "lr": [params_lr[key][measure] for key in params_lr.keys()],
            "x": [key[4:] for key in params_xgb.keys()]}

fig = go.Figure()
fig.add_trace(go.Scatter(x=measures["x"], y=measures["xgb"],
                        mode='lines', name='XGBoost Performance'))
fig.add_trace(go.Scatter(x=measures["x"], y=measures["lr"],
                        mode='lines', name='Logistic Regression Performance'))
fig.update_layout(paper_bgcolor='rgba(0,0,0,0)', plot_bgcolor='rgba(0,0,0,0)',
                  title="{0} - Best Measure for all run".format(measure.upper()))
fig.show()

measures_one_model = {"xgb": [model_first_year_xgb[key][measure] for key in
↳model_first_year_xgb.keys()],
                      "lr": [model_first_year_lr[key][measure] for key in
↳model_first_year_lr.keys()],
                      "x": [key[4:] for key in model_first_year_xgb.keys()]}

fig = go.Figure()
fig.add_trace(go.Scatter(x=measures_one_model["x"], y=measures_one_model["xgb"],
                        mode='lines', name='XGBoost Performance'))
fig.add_trace(go.Scatter(x=measures_one_model["x"], y=measures_one_model["lr"],
                        mode='lines', name='Logistic Regression Performance'))
fig.update_layout(paper_bgcolor='rgba(0,0,0,0)', plot_bgcolor='rgba(0,0,0,0)',
                  title="{0} - Measure for best Model in Year One".format(measure.
↳upper()))
fig.show()

```

```

[16]: with open("../data/cleaned_data/recorded_measures_xgb.json", "r") as f:
        measures_xgb = json.load(f)
    with open("../data/cleaned_data/recorded_measures_lr.json", "r") as f:
        measures_lr = json.load(f)
    xgb_importance = {}
    lr_importance = {}
    for run in measures_xgb:
        for year in run.keys():
            if "importance" in run[year].keys():
                xgb_importance[year] = run[year]["importance"]

    for run in measures_lr:

```

```

for year in run.keys():
    if "importance" in run[year].keys():
        lr_importance[year] = run[year]["importance"]

lr_importance = [val[0] for val in lr_importance.values()]
viz_data = pd.DataFrame(xgb_importance)
name_mapping = {"".join(c for c in col if c.isalnum()):col for col in dfs[0].
    ↳columns}
viz_data = viz_data.rename(name_mapping, axis=0)

year_one_importance = viz_data.loc[np.isfinite(viz_data["XGB Year 1"]), "XGB_
    ↳Year 1"].sort_values(ascending=False)
year_five_importance = viz_data.loc[np.isfinite(viz_data["XGB Year 5"]), "XGB_
    ↳Year 5"].sort_values(ascending=False)

fig = go.Figure()
fig.add_trace(go.Bar(x=year_one_importance.index, y=year_one_importance,
    name='XGB Feature Importance (Gain)'))
fig.update_layout(paper_bgcolor='rgba(0,0,0,0)', plot_bgcolor='rgba(0,0,0,0)',
    ↳title="Year 1 Feature Importances")
fig.show()

fig = go.Figure()
fig.add_trace(go.Bar(x=year_five_importance.index, y=year_five_importance,
    name='XGB Feature Importance (Gain)'))
fig.update_layout(paper_bgcolor='rgba(0,0,0,0)', plot_bgcolor='rgba(0,0,0,0)',
    ↳title="Year 5 Feature Importances")
fig.show()

fig = go.Figure()
fig.add_trace(go.Bar(x=year_one_importance.index, y=year_one_importance,
    name='XGB Feature Importance (Gain) Year 1'))
fig.add_trace(go.Bar(x=year_five_importance.index, y=year_five_importance,
    name='XGB Feature Importance (Gain) Year 5'))
fig.update_layout(paper_bgcolor='rgba(0,0,0,0)', plot_bgcolor='rgba(0,0,0,0)',
    ↳title="Feature Importances",
    legend=dict(
        yanchor="top",
        y=0.99,
        xanchor="left",
        x=0.60
    ))
fig.show()

```

```

intersect_values = np.intersect1d(year_five_importance.index,
    ↪ year_one_importance.index)
len(intersect_values)

```

[16]: 22

```

[20]: ##### Feature Selection by Correlation #####
pd.set_option('chained_assignment', None)
cost_weight = 13
features_path = "../data/cleaned_data/features.json"
dfs, file_names = load_arff_files("../data/raw_data/")
dump_file = False

if os.path.exists(features_path):
    with open(features_path, "r") as f:
        features = json.load(f)
    if "corr_selector_max" not in features:
        features["corr_selector_max"] = {}
    if "corr_selector" not in features:
        features["corr_selector"] = {}
    if "greedy_selector" not in features:
        features["greedy_selector"] = {}
else:
    features = {"corr_selector": {}, "greedy_selector": {}, "corr_selector_max":
    ↪ {}}

for i, df in enumerate(dfs):
    corr_features, corr_measure = main_corr(data=df.copy(),
    ↪ pred_function=pred_function(cost_weight),
                                early_stopping=5, direction="max")

    features["corr_selector_max"]["Year{}".format(i + 1)] = {"features":
    ↪ corr_features, "auc": corr_measure}
    if dump_file:
        with open(features_path, "w") as f:
            json.dump(features, f)
    print(corr_features)

for i, df in enumerate(dfs):
    corr_features, corr_measure = main_corr(data=df.copy(),
    ↪ pred_function=pred_function(cost_weight),
                                early_stopping=10, direction="min")

    features["corr_selector"]["Year{}".format(i + 1)] = {"features":
    ↪ corr_features, "auc": corr_measure}
    if dump_file:

```

```

        with open(features_path, "w") as f:
            json.dump(features, f)
    print(corr_features)

```

Run: 0

D:\CodingProjects\AccentureApplication\src\feature_selecting.py:66:

RuntimeWarning:

invalid value encountered in double_scalars

Chosen column: (gross profit + depreciation) / sales - with Measure value:

0.667866638701911

[0.667866638701911]

Run: 1

KeyboardInterrupt

```

[21]: ##### Greedy Feature Selection ##### CONSUMES A LOT OF TIME!
pd.set_option('chained_assignment', None)
cost_weight = 13
features_path = "../data/cleaned_data/features.json"
dfs, file_names = load_arff_files("../data/raw_data/")

if os.path.exists(features_path):
    with open(features_path, "r") as f:
        features = json.load(f)
    if "corr_selector_max" not in features:
        features["corr_selector_max"] = {}
    if "corr_selector" not in features:
        features["corr_selector"] = {}
    if "greedy_selector" not in features:
        features["greedy_selector"] = {}
else:
    features = {"corr_selector": {}, "greedy_selector": {}, "corr_selector_max":
→ {}}

for i, df in enumerate(dfs):
    if "Year{}".format(i + 1) in features["greedy_selector"]:
        continue
    features_greedy, greedy_measure = main_greedy(data=df.copy(),
→ pred_function=pred_function(cost_weight))
    features["greedy_selector"]["Year{}".format(i + 1)] = {"features":
→ features_greedy, "auc": greedy_measure}
    if dump_file:
        with open(features_path, "w") as f:

```

```
        json.dump(features, f)
    print(features_greedy)
```

Run: 0

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-21-17cab3be30b5> in <module>
    20     if "Year{}".format(i + 1) in features["greedy_selector"]:
    21         continue
--> 22     features_greedy, greedy_measure = main_greedy(data=df.copy(),
    ↪ pred_function=pred_function(cost_weight))
    23     features["greedy_selector"]["Year{}".format(i + 1)] = {"features":
    ↪ features_greedy, "auc": greedy_measure}
    24     if dump_file:

D:\CodingProjects\AccentureApplication\src\__main_fs__.py in main_greedy(data,
    ↪ pred_function, early_stopping, tolerance, verbose, max_processes)
    15     features, measures = gfs.run_selection(data=data, label=label,
    ↪ prediction_function=pred_function,
    16
    ↪ early_stopping_iter=early_stopping,
--> 17                                     tolerance=tolerance,
    ↪ verbose=verbose, max_processes=max_processes)
    18
    19     return features, measures

D:\CodingProjects\AccentureApplication\src\feature_selecting.py in
    ↪ run_selection(self, data, label, prediction_function, early_stopping_iter,
    ↪ tolerance, verbose, max_processes)
    115         if max_processes > 1:
    116             measures = self.calculate_parallel(chosen_columns,
    ↪ remaining_columns, prediction_function, data, label,
--> 117                                     max_processes)
    118         else:
    119             measures = self.calculate_sequential(chosen_columns,
    ↪ remaining_columns, prediction_function, data, label)

D:\CodingProjects\AccentureApplication\src\feature_selecting.py in
    ↪ calculate_parallel(chosen_columns, remaining_columns, prediction_function,
    ↪ data, label, max_processes)
    152         for p in processes:
    153             while True:
--> 154                 if semaphore.acquire(timeout=1):
    155                     break
    156                 p.start()
```

KeyboardInterrupt:

```
[26]: ##### Evaluation Feature Selection #####
with open("../data/cleaned_data/features.json","r") as f:
    features = json.load(f)

measures_year_one = {sel: features[sel]["Year1"]["auc"] for sel in features.
    ↪keys()}

year_one = {sel: features[sel]["Year1"] for sel in features.keys()}

for i in range(1, 6):
    year = "Year{}".format(i)
    print(year)
    for sel in features.keys():
        if year in features[sel]:
            print("Selector : {}, Number Features: {}, Measure: {}".
                ↪format(sel, len(features[sel]["Year{}".
                    ↪format(i)]["features"]),
                    features[sel]["Year{}".format(i)]["auc"][-1]))

print("Union of features between greedy and corr sel: {}".format(
    len(np.intersect1d(year_one["corr_selector"]["features"],
    ↪year_one["greedy_selector"]["features"])))
)
print("Union of features between greedy and corr max sel: {}".format(
    len(np.intersect1d(year_one["corr_selector_max"]["features"],
    ↪year_one["greedy_selector"]["features"])))
)
fig = go.Figure()
x = list(range(70))
fig.add_trace(go.Scatter(x=x, y=measures_year_one["corr_selector"],
    ↪mode='lines', name='Correlation Selector Min_
    ↪Correlation'))
fig.add_trace(go.Scatter(x=x, y=measures_year_one["corr_selector_max"],
    ↪mode='lines', name='Correlation Selector Max_
    ↪Correlation'))
fig.add_trace(go.Scatter(x=x, y=measures_year_one["greedy_selector"],
    ↪mode='lines', name='Greedy Selector'))
fig.update_layout(paper_bgcolor='rgba(0,0,0,0)', plot_bgcolor='rgba(0,0,0,0)',
    ↪title=measure.upper(), legend=dict(
        yanchor="top",
        y=0.29,
        xanchor="left",
        x=0.60
    ))
```



```
fig.show()
```

Year1

Selector : corr_selector, Number Features: 18, Measure: 0.8796136384987296

Selector : greedy_selector, Number Features: 16, Measure: 0.9246464533917096

Selector : corr_selector_max, Number Features: 10, Measure: 0.7184320511055909

Year2

Selector : corr_selector, Number Features: 42, Measure: 0.8652532487465465

Selector : corr_selector_max, Number Features: 6, Measure: 0.7331988642177427

Year3

Selector : corr_selector, Number Features: 12, Measure: 0.8376566423628773

Selector : corr_selector_max, Number Features: 23, Measure: 0.7406769533867856

Year4

Selector : corr_selector, Number Features: 29, Measure: 0.8725941910832826

Selector : corr_selector_max, Number Features: 6, Measure: 0.6934307730466096

Year5

Selector : corr_selector, Number Features: 17, Measure: 0.8661359201773835

Selector : corr_selector_max, Number Features: 13, Measure: 0.8439465631929047

Union of features between greedy and corr sel: 4

Union of features between greedy and corr max sel: 4

[]: